

Project: Fast Fourier Transform (FFT)

BASELINE

Date: Wed Dec 1 11:22:32 2021

Version: 2021.1 (Build 3247384 on Thu Jun 10 19:36:33 MDT 2021)

Project: Project4_FFT_Stages

Solution: solution1 (Vivado IP Flow Target)

Product family: zynq

Target device: xc7z020-clg400-1

Timing Estimate

Target	Estimated	Uncertainty
10.00 ns	7.256 ns	2.70 ns

Performance & Resource Estimates

INFO: [SIM 2] *** CSIM start *******

INFO: [SIM 4] CSIM will launch GCC as the compiler.

Compiling/..../..../Desktop/WES237C_Fall2021/Read_the_docs-master/project_files/project4/FFT/HLS/2_Skeleton_Restructured/fft_test.cpp in debug mode

Compiling/..../..../Desktop/WES237C_Fall2021/Read_the_docs-master/project_files/project4/FFT/HLS/2_Skeleton_Restructured/fft.cpp in debug mode

Generating csim.exe

RMSE(R) RMSE(I)
8742.694335937500000 9995.070312500000000

FAIL: Output DOES NOT match the golden output

ERR: [SIM 100] 'csim_design' failed: nonzero return value.

INFO: [SIM 3] *** CSIM finish *******

Question 1: What changes would this code require if you were to use a custom CORDIC similar to what you designed for Project: CORDIC? Compared to a baseline code with HLS math functions for cos() and sin(), would changing the accuracy of your CORDIC core make the DFT hardware resource usage change? How would it affect the performance? Note that you do not need to implement the CORDIC in your code, we are just asking you to discuss potential tradeoffs that would be possible if you used a CORDIC that you designed instead of the one from Xilinx.

ANS: What would be different to use a custom CORDIC would be having to specify the library search path, we would have more control. However, the main drawback of the conventional CORDIC is that the number of iterations is equal to the number of angle constants. Among a great deal of research to overcome this disadvantage, angle recording method is an effective method because it is capable of reducing maybe 50% of the number of iterations.

the CORDIC algorithm stops converging after a point. This is because the successive approximation step sizes is equal to the magnitude of angle, where j is the iteration number. Custom has a effective limit for the sine and cosine functions' input magnitude. If the value of t is constrained to be smaller than this limit, then the CORDIC is used directly. Otherwise, the magnitude of t must be tested and if it is above the limit, an alternative software algorithm must be used, such as the math.h library functions $\cos()$ and $\sin()$.

Question 2: Rewrite the code to eliminate these math function calls (i.e. $\cos()$ and $\sin()$) by utilizing a table lookup. How does this change the throughput and resource utilization? What happens to the table lookup when you change the size of your DFT?

ANS: High precision arbitrary $\sin()$ and $\cos()$ in particular tend to be expensive operations. The coefficients of \cos and \sin have been calculated in advance and the constant table is easy to read directly, so the code only needs to be accessed through

```
W_real and W_imag. ( W_real[k]=cos(-2*PI*k/SIZE);)
```

Simply getting a number from a list is much faster than calculating the number with an algorithm or using a trigonometric function. The primary disadvantage of lookup tables is their memory usage and that continues as the DFT size is increased.

Question 3: Modify the DFT function interface so that the input and outputs are stored in separate arrays. Modify the testbench to accommodate this change to DFT interface. How does this affect the optimizations that you can perform? How does it change the performance? And how does the resource usage change? You should use this modified interface for the remaining questions.

ANS: The FFT is calculated in two parts. The first one transforms the original data array into a bit-reverse order array by applying the bit-reversal method. This makes the mathematical calculations of the second easier. The second part processes the FFT in $N \cdot \log_2(N)$ operations.

Question 4: Loop Optimizations: Examine the effects of loop unrolling and array partitioning on the performance and resource utilization. What is the relationship between array partitioning and loop unrolling? Does it help to perform one without the other? Plot the performance in terms of number of DFT operations per second (throughput) versus the unroll and array partitioning factor. Plot the same trend for resources (showing LUTs, FFs, DSP blocks, BRAMs). What is the general trend in both cases? Which design would you select? Why?

ANS: A 1024-point FFT of complex 32-bit floating point values, running at 250 MHz would require $1024 \text{ points} \cdot (8 \text{ bytes/point}) \cdot 250 \cdot 10^9 \text{ Hz} = 1 \text{ Terabyte/second}$ of data into the FPGA. The relationship between array partitioning and loop unroll is that it tries to achieve kernel computation optimization by expanding the processing code to match that data path. The logic is

to consume all the data as soon as it arrives at the kernel. I have to use both together for it to compile.

Question 5: Dataflow: Apply dataflow pragma to your design to improve throughput. You may need to change your code and make submodules so that it aligns with the task-level or function-level modularity that dataflow can exploit; Xilinx provides some examples of dataflow code. The HLS User Guide pages 145-154 and this summary provide more information. How much improvement does dataflow provide? How does dataflow affect resource usage? What about BRAM usage specifically? Did you modify the code to make it more amenable to dataflow? If so, how? Please describe your architecture(s) with figures on your report.

ANS: The dataflow directive attempts to best implement the memory used to pass data between the two functions. The second for loop labeled butterfly performs all of the butterfly operations for the current stage. This butterfly for loop has another nested for loop with the label DFTpts; each iteration of this for loop performs one butterfly operation. Remember that we are dealing with complex numbers, thus we must perform complex additions and multiplications. The first line in this DFTpts for loop determines the offset of the butterfly.

#pragma HLS dataflow

```
//Call fft
DTYPE Stage1_R[SIZE], Stage1_I[SIZE];
DTYPE Stage2_R[SIZE], Stage2_I[SIZE];
DTYPE Stage3_R[SIZE], Stage3_I[SIZE];
DTYPE Stage4_R[SIZE], Stage4_I[SIZE];
DTYPE Stage5_R[SIZE], Stage5_I[SIZE];
DTYPE Stage6_R[SIZE], Stage6_I[SIZE];
DTYPE Stage7_R[SIZE], Stage7_I[SIZE];
DTYPE Stage8_R[SIZE], Stage8_I[SIZE];
DTYPE Stage9_R[SIZE], Stage9_I[SIZE];
DTYPE Stage10_R[SIZE], Stage10_I[SIZE];

bit_reverse(X_R, X_I, Stage1_R, Stage1_I);
fft_stage_first(Stage1_R, Stage1_I, Stage2_R, Stage2_I);
fft_stages(Stage2_R, Stage2_I, 2, Stage3_R, Stage3_I);
fft_stages(Stage3_R, Stage3_I, 3, Stage4_R, Stage4_I);
fft_stages(Stage4_R, Stage4_I, 4, Stage5_R, Stage5_I);
fft_stages(Stage5_R, Stage5_I, 5, Stage6_R, Stage6_I);
fft_stages(Stage6_R, Stage6_I, 6, Stage7_R, Stage7_I);
fft_stages(Stage7_R, Stage7_I, 7, Stage8_R, Stage8_I);
fft_stages(Stage8_R, Stage8_I, 8, Stage9_R, Stage9_I);
fft_stages(Stage9_R, Stage9_I, 9, Stage10_R, Stage10_I);
fft_stage_last(Stage10_R, Stage10_I, OUT_R, OUT_I);

}
```

Date: Tue Dec 7 20:43:16 2021

Version: 2021.1 (Build 3247384 on Thu Jun 10 19:36:33 MDT)

Project: Project4_FFT_Stages

Solution: solution1 (Vivado IP Flow Target)

Product family: zynq

Target device: xc7z020-clg400-1

Timing Estimate

Target

Estimated

Uncertainty

10.00 ns

9.591 ns

2.70 ns

Performance & Resource Estimates

Figure 1 wo #DATAFLOW

Most of FFT algorithms decomposes the overall N-point DFT into successively smaller and smaller transforms known as a butterfly. The butterfly is the heart of the FFT. It takes data words from memory and computes the FFT. Results are written back to the same memory locations since an in-place algorithm is used. We modify our code into stages. The FFT contains significant parallelism, because each butterfly is independent of one another in the same stage. The $n/2$ butterfly computations every clock cycle with a task interval of 1. When adding the #dataflow the parallelism in the FFT can be exploited.

Question 6: Best architecture: Briefly describe your “best” architecture. In what way is it the best? What optimizations did you use to obtain this result? What are the tradeoffs that you considered in order to obtain this architecture?

ANS:

The pipeline is implemented, but the resources available on the hardware chip are not utilized to the greatest extent. Through loops, the `#Pragma unroll` parallelism can be further increased, the number of loops can be reduced, and the final latency can be reduced. Here are some of the optimizations have been made or tried:

1. Use template to implement `fft_stages`, so you can enter the stage number as a constant at compile time, which is convenient for pragma writing
2. `bit_reverse` Function I and reversed situation of address classifies discussed, to avoid the occurrence of read and write the same address, while explicitly informed that the write after write HLS (WAW) dependence does not exist, and thus the cycle can be done $II = 1$
3. Reduce the `fft_stage_first` sum `fft_stage_last` to a loop, because there is a layer of loops with a loop count of 1
4. Set `dft_loop` the start of the loop to `0~loop_times`, the number of loops can be determined at compile time, so that HLS can give performance results
5. For the operation `dft_loop` of reading `X_R` and `X_I` summing, create a local variable for storage to avoid repeated reading in the following calculations

- To `dft_loop` perform unroll operations, the corresponding input and output arrays are also cyclic divided

FFT Stages Optimized

Date: Tue Dec 7 22:34:05 2021
Version: 2021.1 (Build 3247384 on Thu Jun 10 19:36:33 MDT)
Project: Project4_FFT_Stages

Solution: solution1 (Vivado IP Flow Target)
Product family: zynq
Target device: xc7z020-clg400-1

Timing Estimate

Target	Estimated	Uncertainty
10.00 ns	7.256 ns	2.70 ns

Performance & Resource Estimates

Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
fft				-	4508	4.508E4	-	2051	-	dataflow	652	1482	195719	287605	0
bit_reverse				-	2050	2.050E4	-	2050	-	no	0	0	108	1430	0
fft_stage_first				-	81	810.000	-	81	-	no	0	144	16346	24454	0
fft_stages_2_s				-	274	2.740E3	-	274	-	no	0	36	5029	7511	0
fft_stages_3_s				-	146	1.460E3	-	146	-	no	0	78	8729	13175	0
fft_stages_4_s				-	82	820.000	-	82	-	no	0	174	19281	28615	0
fft_stages_5_s				-	101	1.010E3	-	101	-	no	0	186	25383	37178	0
fft_stages_6_s II Violation				-	179	1.790E3	-	179	-	no	28	192	26139	38366	0
fft_stages_7_s II Violation				-	261	2.610E3	-	261	-	no	28	192	26132	38365	0
fft_stages_8_s II Violation				-	425	4.250E3	-	425	-	no	28	192	26125	38349	0
fft_stages_9_s II Violation				-	753	7.530E3	-	753	-	no	28	192	26118	38331	0
fft_stage_last II Violation				-	146	1.460E3	-	146	-	no	12	96	16137	19293	0

Generating csim.exe

```

=====
RMSE(R)      RMSE(I)
0.000450454419479 0.000541143584996
=====

```

```

*****

```

PASS: The output matches the golden output!

```

*****

```

1.2) Generate RTL code and export it

Synthesis Summary Report of 'axi4_sqrt'

General Information

Date: Thu Dec 2 08:36:50 2021

Version: 2021.1 (Build 3247384 on Thu Jun 10 19:36:33 MDT 2021)

Project: axi4_sqrt

Solution: solution1 (Vivado IP Flow Target)

Product family: zynq

Target device: xc7z020-clg400-1

Timing Estimate

Target	Estimated	Uncertainty
10.00 ns	7.300 ns	2.70 ns

Performance & Resource Estimates

Modules && Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM	
axi4_sqrt				-	-	-	-	-	0	-	no	10	0	1850	2420	0
axi4_sqrt_Pipeline_1				-	-	-	-	-	-1	-	no	0	0	59	83	0
axi4_sqrt_Pipeline_VITIS_LOOP_16_1				-	146	1.460E3	-	146	-	-	no	0	0	179	102	0
axi4_sqrt_Pipeline_3				-	-	-	-	-	-1	-	no	0	0	163	146	0

HW Interfaces

M_AXI

Interface	Data Width (SW->HW)	Address Width	Latency	Offset	Register	Max Widen Bitwidth	Max Read Burst Length	Max Write Burst Length	Num Read Outstanding	Num Write Outstanding
m_axi_input_r	32 -> 32	64	0	slave	0	0	16	16	16	16
m_axi_output_r	32 -> 32	64	0	slave	0	0	16	16	16	16

Question 7: Streaming Interface Synthesis: Modify your design to allow for streaming inputs and outputs using `hls::stream`. You must write your own testbench to account for the function interface change from `DTYPE` to `hls::stream`. NOTE: your design must pass Co-Simulation (not just C-Simulation). You can learn about `hls::stream` from the HLS Stream Library. An example of code with both `hls::stream` and `dataflow` is available (along with its testbench) [here](#), and another example showing `hls::stream` between functions. Describe the major changes that you made to your code to implement the streaming interface. What benefits does the streaming interface provide? What are the drawbacks?